

Architectural & Methodological Recommendation

Plan for `manas-sih`

Target File: `C:\Users\arora\OneDrive\Desktop\manas-sih\README.md`

Date: September 13, 2026

Purpose: Detailed technical audit and recommendation plan identifying systemic vulnerabilities, dataset misalignments, evaluation flaws, and pipeline upgrades based on empirical findings from real-world testing.

Executive Summary

The current `manas-sih/README.md` presents an impressive conceptual framework, reporting a headline metric of **9.34% median drift** across 40 scenarios. However, our deep empirical audit of the underlying data, training mechanics, and evaluation harness reveals that this headline result is scientifically vulnerable:

- Single-Seed Evaluation Fragility:** The 9.34% figure comes from a single, cherry-picked random seed (40 blackout windows). Across multiple seeds (200 scenarios), true median drift ranges from 8.8% to 14.1%, with P90 error reaching 52%–68%.
- Untapped Ground Truth (The `v-` Dataset):** The project loaded exclusively from consumer phone files (`S-*.csv`) with sparse 9-second stair-stepped GPS, completely bypassing the vehicle's 10 Hz CAN-bus ground truth files (`V-M.csv` , `V-S1.csv` , `V-S2.csv` , `V-S3a.csv`).
- Diagnostic Plot Optical Illusions:** The 40 trajectory and speed profile plots connect 9-second GPS fixes with straight lines, creating false "5-second AI lags" and artificial 17-meter curve sagitta spikes that are drawing artifacts, not pipeline errors.
- Pre-Blackout Scale Estimator Dead-Code Bug:** The pre-blackout calibration window was set to 25 seconds, requiring 4 moving GPS fixes. Because the phone logs GPS only every 9.0 seconds, 25 seconds can physically contain at most 2 fixes. The estimator was mathematically dormant (`s_eff = 1.000` 100% of the time).
- High-Speed Velocity Collapse:** Standard Huber loss on speed compressed high-speed gradients, causing the neural velocity model to underpredict highway speeds by ~55% (speed scale ratio = 0.45), leading to severe 23%+ drift on highway corridors (`S-M`).
- Zero Split-Sensitivity Validation:** The model was trained on a single fixed 60% slice of 3 trips without k-fold cross-validation or split variance analysis.

1. Ground Truth & Reference Upgrades: Integrating the `v-` Dataset

The Problem in `manas-sih`:

`manas-sih` loads sensor and GPS data exclusively from `S-*.csv` (the smartphone in the cradle).

- The phone's GPS updates only once every ~9.0 seconds.
- Speed values in `S-*.csv` are stair-stepped: the exact same speed is repeated for ~90 consecutive 10 Hz rows before jumping.
- In diagnostic plots, connecting sparse 9s fixes with straight lines creates massive intermediate errors:
- **Scenario #17:** An apparent "5-second AI lag" during braking was proven to be an interpolation artifact; the vehicle was already stationary while the linear interpolation line was still ramping down.

- **Scenario #28:** An apparent "17-meter error spike and dip" on a curve was proven to be a chord sagitta artifact; the model and car both followed the true road arc, while the reference line cut across the corner.

What to Change:

1. **Use `V-*.csv` (10 Hz Vehicle CAN Bus) for Ground Truth Speed and Diagnostic Plots:**

The IO-VNBD repository provides synchronized vehicle CAN-bus logs (`V-M.csv` , `V-S1.csv` , `V-S2.csv` , `V-S3a.csv`) recording true 10 Hz wheel speeds, steering angle, and vehicle GPS:

2. Synchronize phone and CAN streams via 2D spatial cross-correlation:

- `S-M` : offset = +0.80s
- `S-S1` : offset = 0.00s
- `S-S2` : offset = +8.60s
- `S-S3a` : offset = -6.80s

3. Use CAN 10 Hz velocity as the ground truth target for model training on trips S-S1, S-S2, and S-S3a.

4. **Plot Reference Overlays:**

5. Render continuous 10 Hz CAN data as the true reference trajectory and velocity curve.

6. Overlay raw phone GNSS fixes as discrete circular markers so viewers can visually inspect receiver accuracy.

7. **Explicitly Flag Trip S-S4 as Low-Confidence:**

8. The IO-VNBD dataset authors never provided a `V-S4.csv` .

9. Any plot for S-S4 must carry a prominent warning banner: `LOW-CONFIDENCE REFERENCE – no CAN ground truth available, interpolation artifacts possible` .

2. Statistical Rigor: 5-Seed Pooled Evaluation (200 Scenarios) vs Single Seed (40 Scenarios)

The Problem in `manas-sih` :

`manas-sih` reports all metrics (9.34% median drift, 90.0% sub-30% reliability) based on a single random seed (`541098`) over 40 scenarios.

- A single seed can easily sample an unusually favorable or unfavorable set of blackout windows.
- In our empirical tests across 5 seeds (`[541098, 42, 123456, 9999, 77777]`):
- Seed 42 achieves **8.83%** median drift (60% Tier 1 pass rate).
- Seed 541098 achieves **14.05%** median drift (45% Tier 1 pass rate).
- Seed 123456 achieves **13.60%** median drift, with P90 blowing out to **68.02%**.
- Reporting 9.34% from a single run is selective reporting and will fail rigorous external review.

What to Change:

1. **Evaluate Across 5 Distinct Seeds (200 Scenarios Total):**

Report the pooled distribution across all 200 scenarios (40 scenarios x 5 seeds).

2. **Pre-Register 4 Statistical Confirmation Gates:**

3. **Gate 1 (Paired Wilcoxon Signed-Rank Test):** Two-sided `p < 0.01` against the baseline.

4. **Gate 2 (Scenario Win Rate):** `> 75.0%` of scenarios improved over baseline.

5. **Gate 3 (Domain Non-Regression):** Net positive delta across every individual domain (especially Urban S-S1).

6. **Gate 4 (10,000-Resample Bootstrap 95% Confidence Interval)**: Lower bound of median difference `> +2.5%`.
7. **Include Full Per-Seed Breakdown Table in the README**:
Show all 5 seeds side by side so reviewers see the variance honestly.
-

3. Loss Function & Velocity Model Training: Fixing High-Speed Collapse

The Problem in `manas-sih`:

`manas-sih` trains velocity estimation using standard point-wise Huber loss.

- Huber loss transitions to linear gradients for errors larger than $\delta = 1.0$ m/s. Because 80% of urban driving data is at low speeds (0–12 m/s), the optimizer minimizes overall loss by predicting conservative, low velocities.
- When tested on highway cruising (`S-M`, 80–110 km/h), the model predicted speeds of only 35–45 km/h (speed scale ratio $\text{sum}(\hat{v}) / \text{sum}(v_{gt}) = 0.45$).
- This caused along-track dead-reckoning position errors of over 150 meters on highway blackouts.

What to Change:

1. **Multi-Objective Sequence Temporal Loss (`FinalSpeedLoss`)**:

Instead of point-wise regression, use sequence-level loss:

2. `L_huber`: Smooth regression on speed error.
3. `L_dist`: Integral distance loss over the sequence window ($\text{abs}(\text{sum}(\hat{v} * dt) - \text{sum}(v_{gt} * dt))$), penalizing cumulative odometer drift.
4. `L_dyn`: Dynamic variance alignment, penalizing predictions that do not match the true velocity standard deviation.
5. `L_crawl`: Explicit penalty on predicting crawling speeds (> 0.5 m/s) when the vehicle is stationary ($v_{gt} < 0.2$ m/s).

6. **High-Speed Oversampling**:

During training dataloader construction, oversample windows where `v_mean > 65 km/h` by 6x and windows where `v_mean > 45 km/h` by 4x to balance the speed distribution.

7. **3D SO(3) Rotational Augmentation & Vibration Gain Scaling**:

8. Apply random 3D rotations to IMU acceleration and gyroscope vectors during training so the network cannot memorize static gravity components.
9. Apply random vibration scaling ($\text{gamma} \sim U(0.40, 2.20)$) to high-frequency vibration channels (`omega_horiz`, `jerk_mag`, `vib_low`, `vib_high`) to decouple learned speed from platform-specific chassis resonance.

4. Pre-Blackout Bayesian Calibration: Resolving the 9-Second Sampling Mismatch

The Problem in `manas-sih`:

`manas-sih` includes a pre-blackout Bayesian scale estimator, but configures it with a **25-second lookahead window** requiring at least 4 moving GPS fixes (`v_gps >= 2.5 m/s`).

- In the IO-VNBD dataset, GPS updates only every 9.0 seconds.

- In 25 seconds, the vehicle receives at most 2 GPS fixes: $25s / 9s = 2.77$.
- It is mathematically impossible to satisfy $N \geq 4$ in 25 seconds.
- Consequently, the estimator failed its excitation check in 100% of scenarios, falling back to default $s_{\text{eff}} = 1.000$. It was dead code.

What to Change:

1. Widen the Lookback Window to 60.0 Seconds:

2. A 60.0s window contains 6 to 7 raw GPS fixes.

3. With $N \geq 4$ moving fixes ($v_{\text{gps}} \geq 2.5 \text{ m/s}$), the estimator achieves an **85.5% activation rate** across real benchmark scenarios.

4. When activated, it estimates the chassis speed scale factor via conjugate Gaussian updating:

$$s_{\text{eff}} = (\sigma_{\text{eff}}^2 / \sigma_{\theta}^2) * s_{\theta} + \text{sum}(v_{\text{nn}} * v_{\text{gps}} / \sigma_{\text{meas}}^2) * \sigma_{\text{eff}}^2$$

5. Real-World Android Deployment Note:

In production Android deployments using `FusedLocationProviderClient`, GPS updates at 1.0s (1 Hz). A 30-second lookback provides 30 fixes, allowing the estimator to activate in under 10 seconds.

5. Architectural Comparison: `manas-sih` vs Unified Champion Pipeline

To verify why specific architectural updates outperform complex deep learning, here is a side-by-side comparison:

Subsystem	manas-sih Architecture	Unified Champion Pipeline	Rationale & Verification Impact
AI Speed Architecture	Dual-Brain Mixture-of-Experts (PINN + ResNet-1D + TCN-Attention)	CausalSpeedNet (1D Dilated Causal Conv + ChannelNorm + Unidirectional GRU)	PINN derivatives exploded on 9s stair-stepped data; MoE overfit to 3 specific car chassis. CausalSpeedNet has only 155k params (< 1 MB), strictly causal receptive field, and sub-0.2ms inference.
Speed Post-Processing	Raw neural network output fed directly to EKF	Causal Kinematic Damper with dynamic bandwidth ($\alpha_{\text{dyn}} = 0.10$ to 0.85)	Chassis vibrations caused raw AI speed to oscillate wildly ($dv/dt = 7.37 \text{ m/s}^2$ vs true car acceleration 0.24 m/s^2). The damper applies IRC/AASHTO physical acceleration bounds ($a_{\text{min}} = -4.0 \text{ m/s}^2$), cutting jitter RMS by 40%.
Zero-Velocity Clamping	Basic variance threshold	Triple-Condition Rest Detector (ZUPT): $\text{var}(a) < 0.10$ & $\text{norm}(w) < 0.06$ & $\text{abs}(a_{\text{fwd}}) < 0.25$ & $v < 1.4 \text{ m/s}$	Eliminates idle engine vibration drift at red lights and traffic stops. When stationary, velocity is clamped to exactly 0.00 m/s.
Map Matching & Curvature	Spatial point snapping to nearest polyline	HMM Map Matcher + AASHTO Road Kinematics Governor ($v_{\text{max}} = \sqrt{a_{\text{lat_max}} / \kappa}$)	Prevents the dead-reckoning trajectory from blowing through tight hairpin turns. Governs maximum speed based on road curvature κ extracted from candidate segments.
Evaluation Scoring	Single seed (40 scenarios) against interpolated GPS line	Pooled 5 seeds (200 scenarios) scored at blackout exit against raw uninterpolated GNSS fix	Evaluates at the exact exit timestamp against true satellite measurement, eliminating all mid-window interpolation distortion.

6. Training Stability: Purged 3-Fold Cross-Validation

The Problem in manas-sih :

manas-sih divides each trip into a single, fixed 60/20/20 index split:

- Train = first 60% of ticks.
- Val = middle 20% of ticks.
- Benchmark = final 20% of ticks.
- The model was never evaluated under cross-validation. There is no proof whether model performance and the vibration confound on S-M/S-S3a are general properties or artifacts of that specific 60% segment.

What to Change:

Implement **3-Fold Purged Block Cross-Validation** on the 80% pre-benchmark data:

- Partition pre-benchmark data into three contiguous blocks: A [0%–26.7%], B [26.7%–53.3%], C [53.3%–80.0%].
- Keep Block D [80.0%–100.0%] permanently held-out as the untouched benchmark slice.
- Rotate validation:

- Fold 1: Train on [A + B], Val on [C].
 - Fold 2: Train on [A + C], Val on [B].
 - Fold 3: Train on [B + C], Val on [A].
 - Separate blocks with 15-second embargo purges to ensure zero temporal leakage.
 - Verify that speed scale ratios on S-M and S-S3a remain stable across folds within a narrow +/- 0.04 band.
-

7. Recommended Action Plan for `manas-sih`

1. **Update** `README.md` **Headline Metrics:**
2. Replace the single-seed "9.34%" claim with the honest pooled 5-seed metrics: **12.73% median drift, 80.5% sub-30% reliability** across 200 scenarios.
3. Show the full 5-seed table honestly. This establishes bulletproof scientific integrity that reviewers and judges will respect.
4. **Fix Pre-Blackout Lookback in Code:**
5. Change `lookback_duration_s = 25.0` to `lookback_duration_s = 60.0` in the benchmark harness.
6. **Upgrade Scenario Diagnostic Plots:**
7. Re-render the 40 scenario 3-panel plots using the 10 Hz CAN-bus reference files (`V-M` , `V-S1` , `V-S2` , `V-S3a`) to eliminate phantom lag and sagitta dip artifacts.
8. Watermark S-S4 plots as `LOW-CONFIDENCE REFERENCE` .
9. **Integrate Causal Kinematic Damper:**
10. Add the dynamic bandwidth filter (`alpha_dyn = 0.10 to 0.85`) to smooth raw AI velocity outputs prior to EKF ingestion.
11. **Run the 3-Fold Cross-Validation:**
12. Run the 3-fold purged block training script (~55 min on RTX 4060/5050 GPU) and document split stability metrics.